

Ceará Wildfire Digital Twin: An Open-Source Agentic AI Platform with LangGraph

Nauber Gois ^{a,1,*}, Alexandre Lobo ^b, Vladia Celia Monteiro Pinheiro ^c, Daniel Valente de Macedo ^d, Adriano Bessa Albuquerque ^a

^aUniversidade de Fortaleza (UNIFOR), Programa de Pós-Graduação em Informática Aplicada, Fortaleza, CE, Brazil

^bUniversidade de São José (USJ), Macau, China

^cInstituto Federal de Educação, Ciência e Tecnologia do Ceará (IFCE), Fortaleza, CE, Brazil

^dSaxion University of Applied Sciences, Deventer, Netherlands

Abstract

Wildfires threaten biodiversity and ecosystem services in Brazil’s semi-arid Caatinga, yet operational monitoring still depends on centralized products with multi-hour latency. We present an open-source ecological data platform that fuses NASA FIRMS and Open-Meteo in a LangGraph pipeline (10 nodes) for ingestion, validation, parallel geospatial/climate analysis, evidence fusion, risk classification, ReAct diagnosis, and three-class alerting for sustainable fire management. On INPE validation for the 2025 dry season (184 days, 24,573 hotspots), alerts reached 84.2% precision and 91.8% recall, with bootstrap 95% intervals for YES-class precision. Under the same protocol in Maranhão and Piauí, recall stayed above 96% at 58–60% precision, motivating per-state calibration. GOES-16 is research-only (grid-scale $F1 \leq 0.05$ vs. INPE). NeKo-PIGNN attains $R^2 = 0.972$ on synthetic fire dynamics; on short real windows, simpler baselines yield lower RMSE. The MIT-licensed system runs on a single EC2 instance (approximately USD 20/month).

Keywords: wildfire monitoring, Caatinga, computational ecology, ecological data science, physics-informed GNN, LangGraph, NASA FIRMS

1. Introduction

Wildfires are among the greatest ecological threats in Brazil, with cascading impacts on biodiversity, carbon stocks, public health, and rural livelihoods [1, 2]. The State of Ceará sits in the northeastern semi-arid zone, where prolonged dry seasons and the Caatinga — a seasonally dry tropical forest with high endemism and fire-sensitive vegetation — make continuous monitoring essential for ecosystem stewardship and civil defense [3, 4]. From an ecological informatics perspective, the challenge is not only detecting hotspots, but assimilating heterogeneous satellite and climate streams into timely, auditable alerts that support sustainable fire management.

Traditional products such as INPE BDQueimadas and NASA FIRMS provide essential ecological observations but leave three operational gaps: 3–6 hour latency between satellite overpass and data availability [5]; no automatic fusion with local climate and municipal context; and interfaces aimed at remote-sensing specialists rather than environmental managers.

*Corresponding author at Universidade de Fortaleza (UNIFOR), Programa de Pós-Graduação em Informática Aplicada. E-mail: naubergois@gmail.com, francisco.gois@cge.ce.gov.br.

Recent advances in large language models and agent frameworks enable auditable pipelines that combine heterogeneous ecological data and produce natural-language explanations. ReAct reasoning [6] and LangGraph state-graph orchestration [7] are particularly suited to multi-source wildfire workflows.

We adopt *Digital Twin* terminology pragmatically: the platform continuously mirrors satellite and weather observations and hosts predictive modules, but it does not yet provide full bidirectional control or physics-coupled simulation at operational scale [8]. Throughout the paper we refer to the deployable artifact as the *monitoring platform*.

The main contributions of this paper are:

1. **An open-source ecological monitoring platform** that integrates NASA FIRMS (VIIRS/MODIS), Open-Meteo, and Nominatim, with optional GOES-16 for research benchmarks, without an institutional database;
2. **A LangGraph-orchestrated agent pipeline** for collection $\rightarrow$ validation $\rightarrow$ parallel geospatial/climate analysis $\rightarrow$ evidence fusion $\rightarrow$ risk classification $\rightarrow$ ReAct diagnosis $\rightarrow$ three-class alerts;
3. **A RAG module** with FAISS and local embeddings for natural-language queries about methodology and use;
4. **Operational validation** on open data (July 2024–June 2026), including dry-season INPE benchmarks for Ceará and neighboring Caatinga/Cerrado transition states.

The remainder of this paper is organized as follows. Section 2 reviews related work; Section 3 describes the architecture; Section 4 formalizes NeKo-PIGNN; Section 5 summarizes deployment; Section 6 reports results; Section 7 discusses ecological implications and limitations; Section 8 concludes.

2. Related Work

We position the platform at the intersection of ecological Digital Twins, satellite fire remote sensing, physics-informed fire dynamics, and LLM agents for environmental decision support.

2.1. Ecological Digital Twins and Wildfire Systems

Digital Twin surveys establish bidirectional mirroring, real-time updates, and predictive capability as core requirements [9, 10, 11, 12]. Earth-system Digital Twins emphasize large spatiotemporal scales and uncertain natural processes [8, 13, 14].

Wildfire-oriented Twin and agent systems include IVSR [15] (multi-sensor agents; not validated for Brazilian biomes), FIRE-VLM [16] (UAV VLM+RL), FIRETWIN [17] (immersive reconstruction), AIMNET [18] (IoT emission monitoring), GraphFire-X [19] (GNN+XGBoost at the WUI), and FireCastNet [20] (global seasonal GNN). Table 1 compares these systems with the present Caatinga-focused, open-source LangGraph platform.

2.2. Satellite Fire Sensing and Brazilian Biomes

MODIS Collection 6 [21] and VIIRS [22] underpin NASA FIRMS [5] and INPE BDQueimadas [23]; Wooster et al. [24] survey active-fire remote sensing. GOES-16 adds 1–10 min geostationary refresh at ~ 2 km [25]. In Brazil, MapBiomass Fogo maps burned area at 30 m [26];

Table 1: Comparison of Digital Twin systems for wildfires.

System	Agents	Open	BR	RAG	Reported metrics
IVSR [15]	Yes	No	No	No	Latency reduction (qual.)
FIRE-VLM [16]	RL+VLM	No	No	No	6× faster detection
FIRETWIN [17]	No	No	No	No	Visual reconstruction
GraphFire-X [19]	GNN+XGB	Yes	No	No	Vulnerability AUC (Eaton)
FireCastNet [20]	GNN	Yes	Partial	No	F1 > GRU/ConvLSTM
AIMNET [18]	IoT ML	No	No	No	Plume detection rate
This work	LangGraph	Yes	Yes	Yes	84.2% P / 91.8% R

biome studies cover fire drivers [27, 1], ML for management [28], and Landsat burned-area mapping including Caatinga [4]. These products remain largely reactive and rarely couple multi-sensor detection with agent reasoning or predictive modules for regional ecological decision support.

2.3. Koopman Operators, PINNs, and Agentic Pipelines

Koopman theory and DMD variants [29, 30, 31, 32] and deep autoencoder Koopman models [33] enable linear multi-step prediction of nonlinear dynamics; PINNs [34, 35] and Rothermel-based spread models [36, 37] encode fire physics. Hybrid NeKo-PIGNN combines both for interpretability and physical consistency. ReAct agents [6], LangGraph orchestration [7], and RAG [38, 39] provide auditable tool use and documentation search; to our knowledge, LangGraph has not previously been applied to operational wildfire monitoring.

None of the surveyed systems combine, in one open-source stack, near-real-time open satellite data, a LangGraph/ReAct pipeline, RAG, and Caatinga-focused validation without an institutional database — the gap addressed here.

3. System Architecture

Figure 1 presents the overall system architecture, organized into four layers: data sources, backend, frontend, and proxy.




Figure 1: Layered architecture of the Ceará wildfire monitoring platform.

3.1. Data Sources

The system consumes four open data sources:

- **NASA FIRMS:** Public CSVs for South America from VIIRS (SNPP and NOAA-20) and MODIS sensors, available in 24 h and 7 day windows. No authentication required. Filtering by Ceará bounding box (lat: -7.85 to -2.78 ; lon: -41.42 to -37.25).
- **Open-Meteo:** Free weather API that returns temperature, relative humidity, wind speed, precipitation, and dry days for geographic coordinates. No API key required.
- **Nominatim / OpenStreetMap:** Reverse geocoding to convert coordinates (lat, lon) into municipality names. Rate limited to approximately 1 request per second.
- **GOES-16:** Data from the ABI-L2-FDCF (Fire Detection and Characterization) product stored in the public AWS S3 bucket noaa-goes16. Processes consecutive detections, FRP, and pixel temperature. This source is optional in production and used primarily for research benchmarks (Section 6).

Figure 2 illustrates the GOES-16 research path: a multi-scale thermal anomaly grid produced by the unsupervised detection module over Ceará.




Figure 2: Example GOES-16 multi-scale thermal anomaly grid (ABI band 13 brightness-temperature residual) produced by the unsupervised detection module. Higher anomaly scores (dark red) mark persistent thermal clusters treated as candidate fire regions in the research benchmark of Section 6.

3.2. Operation Modes

The system operates in two modes:

Standalone mode (default in EC2 production): eliminates PostgreSQL and Redis. Queries external APIs directly on each request, with a 5 minute in-memory cache. Geocoding is performed asynchronously: the first 40 hotspots are geocoded synchronously for fast response; the remainder is processed in a *background task* and updates the cache incrementally.

Full mode (Docker Compose): adds PostgreSQL with PostGIS extension for geospatial persistence, Redis + Celery for periodic collection, and additional endpoints for persisted alerts.

3.3. FastAPI Backend

The backend is implemented in Python 3.12 with FastAPI (version 0.115). Its main responsibilities are:

- Serving REST endpoints for real data: `/api/v1/real/focos`, `/api/v1/real/clima`, and `/api/v1/real/status`;

- Orchestrating the LangGraph pipeline (Section 3.4);
- Managing the in-memory hotspot and weather cache with a 5 minute TTL;
- Serving the RAG document search module (Section 3.6);
- Managing CORS, logging, and health check.

The system uses dependency injection for the LLM (`llm_factory.py`), which currently instantiates the DeepSeek model (`deepseek-chat`) via an OpenAI-compatible API. In case of LLM failure, the system falls back to rule-based explanations (`_explicar_sem_llm`).

3.4. LangGraph Pipeline

The agent pipeline is implemented as a state graph (`StateGraph`) with 10 nodes, as shown in Figure 3.




Figure 3: LangGraph graph of the agent pipeline. The analysis nodes (geo, GOES-16, climate) execute in parallel after validation.

Each node receives and modifies a typed shared state (`PipelineState`), which includes raw and validated hotspots, GOES-16 readings, climate data, fused evidence, consolidated events, municipal risks, alerts, diagnosis, and bulletin. Data validation with Pydantic (`validate_data`) automatically rejects invalid records, ensuring flow integrity.

3.5. *ReAct Agent*

The diagnostic agent (`react_agent.py`) is built with LangChain (version 0.3+) and follows the ReAct pattern [6]. It provides the following tools:

- `buscar_focos_recentes`: query hotspots by municipality, time window, and source;
- `buscar_dados_climaticos`: temperature, humidity, wind, and precipitation;
- `buscar_risco_municipal`: risk index computed per municipality;
- `buscar_dados_goes16`: GOES-16 readings with FRP and persistence;
- `buscar_historico_mapbiomas`: wildfire history by municipality;
- `listar_municipios_criticos`: ranking of the most critical municipalities.

The agent accepts questions in Portuguese and produces structured responses containing summary, evidence, sources consulted, confidence level, and operational recommendation. The reasoning chain (Thought → Action → Observation) is preserved in the response to ensure auditability. Figure 4 summarizes how LangGraph workers, the ReAct diagnostic loop, and the offline RAG research chat interact.

ReAct agents, LangGraph orchestration, and RAG research chat




Figure 4: ReAct agents, LangGraph orchestration, and RAG research chat. Operational alerts flow through parallel analysis nodes and risk classification; documentation queries use a separate FAISS retrieval path.

3.6. RAG Module – Research Chat

The *research chat* module implements RAG [38, 40]. Six knowledge documents (backend/knowledge/01–06) plus a bilingual glossary are split into 900-character *chunks* with 120-character overlap, indexed in FAISS [39] with BAAI/bge-small-en-v1.5 embeddings [41]. Retrieval uses $\text{top}_k = 5$; response generation is performed by DeepSeek citing the document context.

Unlike the operational agent, which queries live hotspots, the research chat responds exclusively

about methodology, architecture, and application usage — functioning as an intelligent manual accessible from any page of the system.

3.7. React Frontend

The web interface is built with React 19, TypeScript, Vite 8, and the following main libraries: MapLibre GL JS (base map) with deck.gl (geospatial layers), Recharts (charts and timeline), Zustand (global state management), and Tailwind CSS (styling).

The available pages are:

- **Dashboard** (/): executive KPIs (total hotspots, active emergencies, critical municipalities, GOES-16 coverage), detection timeline, municipal risk ranking, and recent alerts;
- **Live Map** (/mapa-real): NASA FIRMS hotspots overlaid on the MapLibre map, with clustering, heatmap, severity filters, and per-click explanations;
- **Alerts** (/alertas): full listing with level filters (informational, attention, alert, emergency);
- **Chat** (/chat): conversational interface with the ReAct agent, including suggested questions, evidence visualization, and reasoning chain;
- **Bulletin** (/boletim): consolidated technical report.

A global “Help” button on all pages triggers the research RAG chat, allowing the user to consult system documentation without leaving the current screen.

3.8. Risk Classifier

The municipal risk index is computed from four components:

$$\begin{aligned} \text{index} = & \min(\text{hotspots} \times 8, 40) \\ & + \text{climate_risk} \times 0.4 \\ & + (\text{GOES16_confirmed} \times 15) \\ & + \min(\text{FRP}/100, 10) \end{aligned} \tag{1}$$

Climate risk, in turn, is derived from:

$$\begin{aligned} \text{climate_risk} = & \max(0, (50 - \text{humidity}) \times 0.4) \\ & + \min(\text{wind} \times 1.5, 20) \\ & + \min(\text{days_without_rain} \times 1.5, 30) \end{aligned} \tag{2}$$

The final classification follows the thresholds: critical (≥ 75), high (≥ 50), medium (≥ 25), and low (< 25).

3.9. Software Stack

The system is implemented in Python 3.12 with FastAPI (backend), React 19 with TypeScript and Vite 8 (frontend), and agent orchestration via LangGraph 0.3. Table 2 summarizes the main dependencies.

Table 2: System software stack.

Layer	Technology	Version
Backend	Python / FastAPI / Uvicorn	3.12 / 0.115 / 0.34
Agents	LangGraph / LangChain	0.3 / 0.3+
LLM	DeepSeek Chat API	deepseek-chat
Embeddings	BAAI/bge-small-en-v1.5 (FAISS)	1.5
Frontend	React / TypeScript / Vite	19 / 6 / 8
Map	MapLibre GL JS + deck.gl	4.7 / 9.1
Styling	Tailwind CSS	4
Proxy	nginx	1.26
Deploy	EC2 Amazon Linux 2023	t2.medium

4. Methodology: Neural Koopman Operator and Physics-Informed GNN

The methodological core of this work combines two complementary paradigms: the *Neural Koopman Operator* for modeling the spatiotemporal dynamics of wildfire hotspots and *Physics-Informed Graph Neural Networks* (PI-GNN) for risk prediction with physics-based regularization grounded in the Rothermel model [36]. This section formally describes both components.

4.1. Neural Koopman Operator for Wildfire Dynamics

Koopman theory [29, 42, 30] establishes that nonlinear dynamical systems can be represented by linear operators in an infinite-dimensional observable space. Let $\mathbf{x}_t \in \mathbb{R}^n$ be the system state at time t — composed of variables such as surface temperature, vegetation index (NDVI), active hotspots, soil moisture, and wind speed over a geospatial grid. The Koopman operator $\mathcal{K}$ acts on observable functions $\psi : \mathbb{R}^n \rightarrow \mathbb{R}$ such that:

$$(\mathcal{K}\psi)(\mathbf{x}_t) = \psi(\mathbf{f}(\mathbf{x}_t)) = \psi(\mathbf{x}_{t+1}) \quad (3)$$

where $\mathbf{f}$ is the underlying nonlinear dynamics of the system. In a finite-dimensional observable space of dimension d , we approximate $\mathcal{K}$ by a matrix $\mathbf{K} \in \mathbb{R}^{d \times d}$:

$$\psi(\mathbf{x}_{t+1}) \approx \mathbf{K}^\top \psi(\mathbf{x}_t) \quad (4)$$

We implement this operator as an encoder-decoder neural network (KoopmanAE) [33, 43]. The encoder $g_\phi : \mathbb{R}^n \rightarrow \mathbb{R}^d$ projects the observed state into latent observables $\mathbf{z}_t = g_\phi(\mathbf{x}_t)$. The linear evolution in latent space is:

$$\mathbf{z}_{t+1} = \mathbf{K}_\theta \mathbf{z}_t \quad (5)$$

where $\mathbf{K}_\theta \in \mathbb{R}^{d \times d}$ is the learned matrix. The decoder $h_\psi : \mathbb{R}^d \rightarrow \mathbb{R}^n$ reconstructs the state from the observables: $\hat{\mathbf{x}}_t = h_\psi(\mathbf{z}_t)$. The KoopmanAE loss function combines three terms:

$$\begin{aligned}
\mathcal{L}_{\text{koop}} = & \underbrace{\|\mathbf{x}_t - h_\psi(g_\phi(\mathbf{x}_t))\|^2}_{\text{reconstruction}} \\
& + \underbrace{\|\mathbf{x}_{t+1} - h_\psi(\mathbf{K}_\theta g_\phi(\mathbf{x}_t))\|^2}_{\text{one-step prediction}} \\
& + \underbrace{\sum_{k=1}^K \|\mathbf{x}_{t+k} - h_\psi(\mathbf{K}_\theta^k g_\phi(\mathbf{x}_t))\|^2}_{\text{multi-step prediction}}
\end{aligned} \tag{6}$$

The classical estimator for $\mathbf{K}$ in a fixed dictionary of observables is Extended Dynamic Mode Decomposition (EDMD) [44], which minimizes the one-step prediction error in observable space,

$$\mathbf{K} = \arg \min_{\mathbf{K}'} \mathbb{E}_t \left[\|\psi(\mathbf{x}_{t+1}) - \mathbf{K}'^\top \psi(\mathbf{x}_t)\|^2 \right], \tag{7}$$

with the closed-form least-squares solution

$$\mathbf{K} = \mathbf{G}^+ \mathbf{A}, \quad \mathbf{G} = \frac{1}{T} \sum_{t=1}^T \psi(\mathbf{x}_t) \psi(\mathbf{x}_t)^\top, \quad \mathbf{A} = \frac{1}{T} \sum_{t=1}^T \psi(\mathbf{x}_t) \psi(\mathbf{x}_{t+1})^\top, \tag{8}$$

where $\mathbf{G}^+$ denotes the Moore–Penrose pseudoinverse. Instead of fixing the dictionary and solving Eq. (8) in closed form, we learn the observables and the operator jointly: $\mathbf{K}_\theta$ is a trainable matrix, initialized near the identity for stability, and optimized end-to-end by gradient descent together with g_ϕ and h_ψ under the loss of Eq. (6) and the spectral penalty of Section 4.4.

The dimensionality d of the observable space is a critical hyperparameter. We use $d = 64$ with temporal cross-validation, balancing reconstruction fidelity (mean error < 0.05) and prediction stability over windows of up to 6 steps (12 hours with 2-hour intervals). Spectral decomposition of the trained $\mathbf{K}_\theta$ shows that $\sim 95\%$ of the dynamic energy concentrates in the 12 dominant coherent modes, supporting this choice of d .

4.2. Physics-Informed Graph Neural Networks

To model the spatial propagation of wildfire risk, we construct a graph $\mathcal{G} = (\mathcal{V}, \mathcal{E}, \mathbf{A})$ where vertices $\mathcal{V}$ represent 1 km^2 grid cells over the territory of Ceará (approximately 149,000 cells), and edges $\mathcal{E}$ connect neighboring cells with weights based on geodesic distance, vegetation type, and prevailing wind direction. The adjacency matrix $\mathbf{A} \in \mathbb{R}^{|\mathcal{V}| \times |\mathcal{V}|}$ is sparse with neighborhood radius $r = 3$ cells.

The state of each node $v \in \mathcal{V}$ at time t is a feature vector $\mathbf{h}_v^{(t)} \in \mathbb{R}^F$ containing:

- Surface temperature (GOES-16 ABI band 14, $11.2 \mu\text{m}$);
- Reflectance (VIIRS bands I1–I5);
- NDVI calculated from MODIS;
- Relative humidity and wind speed (Open-Meteo);

- Hotspot history over the last 7 days;
- Vegetation type (MapBiomas);
- Terrain slope and aspect (SRTM 30m).

Node state updates follow the *message-passing* paradigm with L layers:

$$\mathbf{m}_v^{(l)} = \sum_{u \in \mathcal{N}(v)} \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)} \cdot \mathbf{A}_{vu} \quad (9)$$

$$\mathbf{h}_v^{(l)} = \sigma(\mathbf{m}_v^{(l)} + \mathbf{b}^{(l)}) \quad (10)$$

where $\mathcal{N}(v)$ are the neighbors of v , $\mathbf{W}^{(l)} \in \mathbb{R}^{F \times F}$ and $\mathbf{b}^{(l)} \in \mathbb{R}^F$ are learned parameters, and σ is a ReLU activation. We use $L = 3$ layers.

4.2.1. Physics-Informed Loss with Rothermel

The main innovation is regularizing the GNN loss with the Rothermel fire propagation physics equation [36], following the PINN paradigm [34] extended by scientific machine learning models [35, 45, 46]. The rate of spread R (m/min) is modeled as:

$$R = \frac{I_R \xi (1 + \phi_w + \phi_s)}{\rho_b \epsilon Q_{ig}} \quad (11)$$

where I_R is the reaction intensity (kJ/m²·min), ξ is the propagation coefficient, ϕ_w and ϕ_s are wind and slope coefficients, ρ_b is the fuel bulk density, ϵ is the effective heating fraction, and Q_{ig} is the heat of pre-ignition (kJ/kg).

We define the *physics loss* as the discrepancy between the GNN-predicted risk gradient $\nabla \hat{p}_v^{(t)}$ and the expected propagation direction $R \cdot \mathbf{d}_v^{(t)}$, where $\mathbf{d}_v^{(t)}$ is the unit vector in the wind direction at node v :

$$\mathcal{L}_{\text{phys}} = \frac{1}{|\mathcal{V}|} \sum_{v \in \mathcal{V}} \|\nabla \hat{p}_v^{(t)} - R_v^{(t)} \cdot \mathbf{d}_v^{(t)}\|^2 \quad (12)$$

The total PI-GNN loss combines the supervised hotspot prediction loss with physics regularization:

$$\begin{aligned}
\mathcal{L}_{\text{total}} = & \underbrace{\frac{1}{N} \sum_{i=1}^N \text{BCE}(y_i, \hat{p}_i)}_{\text{classification}} \\
& + \lambda_1 \underbrace{\frac{1}{N} \sum_{i=1}^N \|\mathbf{h}_i - \hat{\mathbf{h}}_i\|^2}_{\text{auto-encoding}} \\
& + \lambda_2 \underbrace{\mathcal{L}_{\text{phys}}}_{\text{physics regularization}} \\
& + \lambda_3 \underbrace{\|\mathbf{W}\|_2}_{\text{L2 regularization}}
\end{aligned} \tag{13}$$

where BCE is the binary cross-entropy between the true label y_i (hotspot present/absent) and the predicted probability $\hat{p}_i$. The hyperparameters $\lambda_1 = 0.1$, $\lambda_2 = 0.05$, and $\lambda_3 = 10^{-4}$ were tuned via validation.

4.3. Koopman-PI-GNN Integration in the Pipeline

Figure 5 illustrates the hybrid NeKo-PIGNN architecture. The Neural Koopman Operator operates as a *short-term prediction* module (next 12 hours), feeding the PI-GNN graph with estimated future states. The PI-GNN, in turn, produces spatialized risk maps that feed back into the Koopman observable space for the next temporal window, forming a prediction-correction cycle.




Figure 5: Hybrid Neural Koopman Operator and PI-GNN architecture. The deterministic KoopmanAE encoder maps the state vector into Koopman observables, which are propagated linearly by the low-rank matrix $\mathbf{K}$ and spatially regularized by the PI-GNN with the Rothermel physics loss; training follows the two-phase curriculum of Section 4.5.

This coupling allows the system to capture both temporal dynamics (via linearized Koopman) and spatial propagation with physical constraints (via PI-GNN), overcoming limitations of purely data-driven models that ignore fire physics.

4.4. Implementation Hyperparameters and Operational Integration

The KoopmanAE encoder g_ϕ is an MLP with 3 hidden layers (512, 256, 128 neurons) and ReLU activation, mapping the state vector $\mathbf{x}_t \in \mathbb{R}^{140}$ (7 variables $\times$ 20 sampled grid cells) to latent observables $\mathbf{z}_t \in \mathbb{R}^{64}$. The Koopman matrix $\mathbf{K}_\theta \in \mathbb{R}^{64 \times 64}$ is trained with spectral stability regularization:

$$\mathcal{L}_{\text{spec}} = \max(0, \rho(\mathbf{K}_\theta) - 1 + \epsilon) \quad (14)$$

where $\rho(\mathbf{K}_\theta)$ is the spectral radius and $\epsilon = 0.01$ ensures a stability margin. Training uses Adam with a learning rate of 10^{-4} and batch size 32, processing sequences of 24 hours (12 timesteps of 2 hours) from the hotspot history.

The GNN is implemented with PyTorch Geometric [47] using GCNConv graph convolutions. Each of the $L = 3$ layers has $F = 128$ hidden channels. The graph is built dynamically at each collection cycle, connecting active grid cells (with at least one hotspot in the last 7 days) to their neighbors within radius $r = 3$ cells (3 km). The physics loss $\mathcal{L}_{\text{phys}}$ is computed offline using Rothermel parameters precomputed for each MapBiomass vegetation type (12 classes) and current climate conditions; the gradient $\nabla \hat{p}_v^{(t)}$ is approximated by finite differences between neighboring graph nodes.

The production LangGraph pipeline (`backend/app/agents/langgraph_pipeline.py`) contains **10 nodes**: `coletar_dados`, `validar_dados`, `agente_geoespacial`, `agente_goes16`, `agente_climatico`, `fundir_evidencias`, `classificar_risco`, `agente_react_diagnostico`, `gerar_alertas`, and `gerar_boletim`. NeKo-PIGNN is *not* embedded as additional graph nodes; trained KoopmanAE and PI-GNN models run as batch inference jobs (`backend/experiments/`) whose outputs feed the municipal risk classifier and three-class alert experiments (Section 6.15). When invoked offline, NeKo-PIGNN adds approximately 2.3 s per batch (Koopman: 0.8 s; PI-GNN: 1.2 s; fusion: 0.3 s) on CPU (Intel Xeon Platinum 8375C, AWS EC2 t2.medium).

4.5. Two-Phase Curriculum Training

Joint end-to-end training of the Koopman operator and the GNN is unstable under data scarcity: the spatial message-passing gradients dominate before the latent dynamics are well-formed, which is the failure mode observed in the Exp-B v1 iteration (Section 6.7). The v2 models are therefore trained with a two-phase curriculum [48] (`backend/experiments/validate_models_v2.py`):

1. **Phase 1 — Koopman warm-up (100 epochs).** Only the encoder g_ϕ , decoder h_ψ , and operator $\mathbf{K}_\theta$ are optimized (AdamW, learning rate 10^{-3} , cosine annealing, gradient-norm clipping at 1.0) under a geometrically damped multi-step objective,

$$\mathcal{L}_{\text{ph1}} = \mathcal{L}_{\text{recon}} + 2\mathcal{L}_{\text{pred}}^{(1)} + \frac{1}{2} \sum_{s=1}^3 2^{-(s-1)} \mathcal{L}_{\text{pred}}^{(s)}, \quad (15)$$

where $\mathcal{L}_{\text{pred}}^{(s)}$ is the s -step prediction error obtained by iterating $\mathbf{K}_\theta$ in latent space. Longer horizons receive exponentially smaller weights, growing the effective prediction horizon as training progresses.

2. **Phase 2 — full model with physics ramp-up (150 epochs).** The GNN layers and output head are activated at half the initial learning rate, and the physics-regularization weight is ramped linearly, $\lambda_{\text{phys}}(\tau) = \min(0.5, 0.5 \tau / \tau_{\text{max}})$, so that Rothermel-consistency constraints are imposed only after the latent dynamics stabilize.

Ablations confirm the curriculum is necessary: without it, KL- or physics-dominated gradients reproduce the v1 collapse ($R^2 \approx 0.28$ – 0.34 vs. 0.972 with curriculum; Tables A.18 and 9).

4.6. Error Decomposition and Spectral Analysis

The k -step prediction error of the hybrid model decomposes into an approximation term (finite-dimensional parameterization of $\mathcal{K}$) and a truncation term (restriction to the d -dimensional observable subspace):

$$\|\mathbf{x}_{t+k} - \hat{\mathbf{x}}_{t+k}\| \leq \underbrace{\|h_\psi(\mathbf{K}_\theta^k \mathbf{z}_t) - \mathcal{K}^k \psi(\mathbf{x}_t)\|}_{\text{approximation error}} + \underbrace{\|\mathcal{K}^k \psi(\mathbf{x}_t) - \psi(\mathbf{x}_{t+k})\|}_{\text{truncation error}}. \quad (16)$$

Under the spectral-stability constraint of Eq. (14), all eigenvalues satisfy $|\lambda_i| \leq 1 - \epsilon$, so the approximation term decays as $\mathcal{O}(e^{-\gamma k})$ for the damped modes, while the truncation term is bounded by the subspace dimension d . The stationary mean prediction error is accordingly bounded by

$$\lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T \|\mathbf{x}_{t+1} - \hat{\mathbf{x}}_{t+1}\|^2 \leq \frac{\sigma_{\text{obs}}^2}{T} + \varepsilon_d, \quad (17)$$

where σ_{obs}^2 is the observational noise variance and ε_d the subspace approximation error. Empirically, the concentration of $\sim 95\%$ of the dynamic energy in 12 coherent modes (Section 4.1) indicates that ε_d is small at $d = 64$, and explains why prediction quality degrades gracefully rather than abruptly over the 6-step (12-hour) operational horizon.

5. Deployment

The system was deployed on an Amazon Linux 2023 EC2 instance (t2.medium — 2 vCPUs, 4 GB RAM). Automated deployment via shell script (`deploy/uniform-deploy.sh`) installs nginx, Python 3.12, Node.js 20, and configures the systemd service `uniform-backend`. The frontend build is served statically by nginx, which reverse proxies `/api/` to unicorn on port 8000.

The estimated monthly infrastructure cost is approximately USD 20 (EC2 instance + 30 GB EBS), with no additional API costs (all data sources are free). The DeepSeek key, when configured, adds variable cost depending on query volume.

6. Experimental Results

This section reports quantitative results integrating four experimental tracks (A–D) documented in `docs/experimentos/` and `backend/experiments/results/`: (i) the operational pipeline, deployed continuously for 24 months (Jul/2024–Jun/2026); (ii) GOES-16 unsupervised benchmark; (iii) Exp-B municipal prediction iterations (v1–v9, synthetic and real); (iv) agent, RAG, and risk-index calibration. Reference INPE BDQueimadas records span January 2024 to June 2026.

6.1. Experimental Design and Protocol

We structured the evaluation program into four complementary tracks, fully documented under `docs/experimentos/` and `backend/experiments/results/`.

Track A — Operational pipeline (24 months, Jul/2024–Jun/2026). Continuous ingestion of NASA FIRMS CSVs, Open-Meteo weather, Nominatim geocoding, LangGraph orchestration (10 nodes), ReAct diagnostics, and RAG queries. Metrics: time-to-first-response, geocoding rate, agent success rate, backend availability, and inter-day FIRMS variability (e.g., 426% day-over-day change on 2024-06-05 vs. 2024-06-06, validating the need for autonomous polling).

Track B — Municipal fire prediction (Exp-B, v1–v9). Iterative benchmark on synthetic then real data (Mar–Jun/2026): (v1–v2) next-day regression on simulated Ceará dynamics; (v3–v4) real FIRMS/INPE/Open-Meteo regression; (v5–v7) rare-event binary detection with Average Precision, persistence priors, and precision modes; (v8–v9) three-class alert system (NO/UNCERTAIN/YES). Temporal split 70/10/20% (67/10/20 days). Models: MLP, LSTM, XGBoost, Koopman-Det, NeKo-PIGNN. Intermediate iteration tables (v1, v4–v7) are in Appendix A.

Track C — GOES-16 unsupervised benchmark. Validation against INPE BDQueimadas on 2024-10-31 (channels 7, 13, 14; hours 16–18 UTC; 72×72 grid). Methods: digital twin assimilation, isolation forest, spatial residual, *combined persistence* (Lines B in `docs/METODOLOGIA_NOVA_PROPOSTA.md`), and *consensus ensemble* (Line E, 2-of-3 vote). Reproducible via `src/unsupervised_fire_goes.py`; event-centered evaluation via `src/pixel_inpe_eval.py`.

Track D — Agent, RAG, and risk-index calibration. 100 operational ReAct questions; 30 RAG architecture queries (FAISS recall@5, manual completeness); heuristic risk classifier calibrated on 111,054 INPE records (docs/calibracao/).

6.2. Reference Datasets and Descriptive Statistics

Table 3 summarizes INPE BDQueimadas records for Ceará (Jan/2024–Jun/2026): 111,639 records across 184 municipalities (99.5% of the state) in the released dataset (data/inpe_focos_ce/); risk-calibration experiments (Section 6.6) used a 111,054-record snapshot taken on 12/Jun/2026, before the final records for May–June 2026 were appended. The dry season (Jul–Dec) concentrates approximately 93% of annual activity in both complete years. Note that 2025 records exceed 2024 by more than an order of magnitude (94,135 vs. 7,160): 2025 combined an extreme dry season with the multi-satellite BDQueimadas export (NOAA-20/21, NPP, AQUA, GOES-19), whereas the 2024 export in the released dataset lacks per-satellite attribution and likely undercounts total detections; year-over-year comparisons should therefore be interpreted with caution.

Figure 6 shows monthly and daily temporal patterns used for alert calibration. The pronounced dry-season peaks justify season-aware thresholding in the three-class alert system (Section 6.10).

Table 3: INPE hotspot descriptive statistics for Ceará, computed from the released dataset (Jan/2024–Jun/2026; 111,639 records).

Metric	2024	2025	2026 (Jan–Jun)
Total hotspots	7,160	94,135	10,344
Mean monthly (dry season)	1,119	14,563	—
Mean monthly (rainy season)	73	1,125	1,724
Affected municipalities (monthly mean)	68	138	97
Peak daily count	464	1,831	854
Days with fire (> 0 hotspots)	198	359	157




Figure 6: Wildfire hotspot temporal patterns in Ceará from INPE BDQueimadas records (Table 3). (a) Monthly totals for 2024–2025, showing the dry-season (Jul–Dec) concentration of $\sim 93\%$ of annual activity; (b) daily time series, whose pronounced peaks (up to 1,831 hotspots/day in 2025) motivate the season-aware thresholds of the three-class alert system (Section 6.10).

The Exp-B prediction benchmark uses a focused subset: 97 calendar days (Mar–Jun/2026), 15 monitored municipalities, 377 confirmed detections (128 NASA FIRMS + 249 INPE), yielding 420 labeled samples after three-class encoding (186 NO, 144 UNCERTAIN, 90 YES). Split: 67 train / 10 validation / 20 test days (70/10/20%). Table 4 summarizes sensor contribution and input features for this subset.

Figure 7 visualizes spatial distribution, temporal evolution, climate drivers, and model comparison on the real-data window. Municipal-level counts are reported in Appendix A (Figure A.11).

Table 4: Exp-B real-data subset: detections by sensor and model features.

(a) Detections by sensor			(b) Input features (v3+)	
Sensor	Count	Res.	#	Feature (source)
VIIRS SNPP	34	375 m	0–4	Climate vars (Open-Meteo)
VIIRS NOAA-20	79	375 m	5	fire_count (FIRMS+INPE)
MODIS	15	1 km	<i>Note: v5+ adds 20-dim lags, neighbors, drought.</i>	
INPE	249	~1 km		
Total	377	—		

6.3. Operational Pipeline Performance


Table 5 summarizes 540 NASA FIRMS collections (Jul/2024–Jun/2026, 38,400+ detections): VIIRS SNPP contributes 48 %, VIIRS NOAA-20 36 %, and MODIS 16 %. First response time averages 52.4 s; the 5-minute cache reduces subsequent requests to under 1 s. Section 6.12 positions these metrics against NASA FIRMS and INPE baselines.

Table 5: Operational hotspot detection and sensor contribution (540 NASA FIRMS collections, Ceará).


(a) Pipeline metrics			(b) Sensor share		
Metric	Value	Unit	Sensor	Hotspots	Share
Mean hotspots / collection	42.3	count	VIIRS SNPP	18,432	48 %
Std. deviation	18.7	count	VIIRS NOAA-20	13,824	36 %
Max. in 24 h	187	count	MODIS	6,144	16 %
Municipalities (mean)	14	count			
Critical/high severity	23 %	—			
Geocoding rate	93 %	—			
Cold-start latency	52.4	s			
Agent success rate	97 %	—			
FIRMS HTTP 503 rate	0.74 %	—			

6.4. Agent and RAG Evaluation


Table 6 consolidates Track D AI component benchmarks. The ReAct agent achieved 97% success on 100 operational questions (8.3 s mean latency); failures were DeepSeek timeouts




(a) Spatial map



(b) Daily series



(c) Climate (Crateús)



(d) Model comparison

Figure 7: Exp-B real-data subset (97 days, Mar–Jun/2026; 377 confirmed detections in 15 monitored municipalities). (a) Spatial distribution of NASA FIRMS and INPE hotspots; (b) daily detection counts over the study window; (c) co-evolution of temperature, humidity, and precipitation for Crateús, the most active municipality; (d) next-day fire-count regression compared across the five benchmarked models (Table 10).

only, with 100% rule-based fallback coverage. RAG recall@5 reached 91% on the original 30-question manual protocol. A formal 50-question retrieval benchmark (Exp-ROBUST-004) yields recall@5 = 98% and MRR = 0.91 after adding a bilingual English–Portuguese glossary chunk (00_english_glossary.md) to the knowledge base; without this index, automated recall@5 was 74%. Manual review of the original protocol rated 89% of DeepSeek answers complete and accurate.

Table 6: AI agent, RAG, and LangGraph performance metrics.

(a) Agent & pipeline			(b) RAG benchmark (50 Q)	
Component	Metric	Value	Metric	Value
ReAct agent	Success rate	97 %	Questions	50
ReAct agent	Mean latency	8.3 s	Indexed chunks	32
ReAct agent	Timeout rate	3 %	Recall@5	98.0%
ReAct agent	Fallback coverage	100 %	MRR	0.907
LangGraph	Full cycle (10 nodes)	2.3 s	Manual accuracy	89 %
KoopmanAE	Inference	0.8 s		
PI-GNN	Inference	1.2 s		

6.5. GOES-16 Unsupervised Detection Benchmark

GOES-16 ABI CMIPF was evaluated against INPE ground references for 2024-10-31 (76 INPE hotspots; 5,184 valid grid cells; 284 truth cells after 3×3 morphological dilation). Low F1 (≤ 0.05) reflects structural misalignments documented in docs/METODOLOGIA_NOVA_PROPPOSTA.md: temporal window mismatch, thermal anomaly vs. active-fire semantics, 2 km GOES cells vs. point-like INPE records, and CMIPF brightness temperature vs. dedicated AF products.

Table 7 contrasts grid-based and pixel-level evaluation modes; Figure 8 visualizes the F1 and recall of each unsupervised method against the INPE reference. Combined persistence and consensus ensemble (Lines B/E) yield zero detections but also zero false positives—suitable for high-stakes alert filtering when fused with FIRMS in the operational pipeline.




Figure 8: GOES-16 unsupervised detection benchmark against INPE BDQueimadas ground truth (2024-10-31, 76 INPE hotspots, grid mode with dilated truth): F1-score and recall per method (digital twin assimilation, isolation forest, persistence, spatial residual, consensus ensemble). All methods score $F1 \leq 0.05$, motivating the FIRMS-based operational design (Table 7).

Table 7: GOES-16 vs. INPE validation (2024-10-31). Grid: dilated truth (284 cells); pixel: 15 km focal matching.

Mode	Method	TP	FP	FN	P	R	F1
Grid	Digital Twin (3h)	30	903	254	0.032	0.106	0.049
Grid	Isolation Forest	18	680	266	0.026	0.063	0.037
Grid	Spatial Residual	8	718	41	0.011	0.163	0.021
Grid	Combined Persistence	0	0	284	0.000	0.000	0.000
Pixel	Cluster match	4	152	72	0.026	0.053	0.034

6.6. Risk Index Calibration and Sensitivity (Track D)

A component-wise sensitivity analysis of the heuristic risk index (Section 3.8) shows that the classifier is dominated by the hotspot-count term (capped at 40 of the ~ 100 attainable points) and by GOES-16 confirmation (+15 points), while the FRP term contributes marginally (capped at 10 points); within the climate component, days without rain and low humidity carry roughly three times the weight of wind speed at typical dry-season values. Severity classifications are therefore most sensitive to hotspot-count and drought-related inputs, which motivates the calibration below.

The heuristic risk classifier was calibrated against a 111,054-record INPE BDQueimadas snapshot (12/Jun/2026; 19,249 positive / 11,362 negative municipal-day samples, Jan/2024–May/2026). Labels are defined directly from the reference data: a municipal-day is *positive* if it contains at least one INPE hotspot and *negative* otherwise. Because the risk index takes the daily hotspot count as its dominant input (Eq. 1), a municipal-day can only exceed the severity threshold when hotspots are present, so precision equals 1.000 *by construction* (false positives are structurally impossible) and

must not be read as an independent validation result. The informative quantities here are recall and F1 — how many fire days the index flags at the calibrated thresholds — and the calibration accordingly maximizes F1. Table 8 reports baseline vs. optimized weights: optimized weights increase hotspot sensitivity (w_{focos} 8→12), drought penalty ($w_{\text{clima_seca}}$ 1.5→2.5), and hotspot cap (max_focos 40→50), with severity thresholds lowered by 5 points. Out-of-sample evaluation against labels independent of the index inputs (e.g., burned-area products such as MapBiomas Fogo) remains future work.

Table 8: Risk index calibration on the INPE BDQueimadas snapshot (111,054 records, 12/Jun/2026). Precision is 1.000 by construction (see text); recall and F1 are the informative metrics.

Metric	Baseline	Optimized	Δ
F1	0.833	0.936	+0.104
Precision	1.000	1.000	—
Recall	0.713	0.880	+0.167
Accuracy	0.820	0.925	+0.105

6.7. Exp-B Iteration History: Synthetic and Real Data

v1 (200 timesteps). Initial Koopman-VAE and NeKo-PIGNN underperformed shallow baselines ($R^2 \approx 0.28$ – 0.34 vs. MLP 0.93) due to KL-regularized VAE competing with prediction (Table A.18, Appendix). **v2 (500 timesteps).** Architecture fixes—deterministic Koopman, curriculum learning, spectral regularization—reversed the gap. Table 9 shows NeKo-PIGNN v2 reaching $R^2=0.972$ on synthetic next-day prediction, establishing an upper bound before real-data noise.

Table 9: Exp-B v2: next-day prediction on synthetic data (500 steps, 30 nodes). Best in **bold**.

Model	RMSE↓	MAE↓	R^2 ↑	F1↑
MLP	0.069	0.024	0.968	0.975
LSTM	0.084	0.045	0.951	0.945
XGBoost	0.087	0.039	0.948	0.931
Koopman-Det (ours)	0.068	0.022	0.968	0.975
NeKo-PIGNN v2 (ours)	0.064	0.024	0.972	0.975

6.8. Exp-B Real-Data Regression (v3–v4)

v3 validated all models on real FIRMS/INPE/Open-Meteo data. MLP led in RMSE/ R^2 ; all models maintained binary F1 > 0.92 on next-day fire-count regression—a *different task* from three-class alert metrics (Section 6.10). The synthetic-to-real R^2 drop (Table 10b) ranges from 17 to 25 percentage points, expected under label noise and data scarcity with only 67 training days.

v4 tested feature engineering with 62–67 training samples; MLP remained best ($R^2=0.782$), while feature expansion degraded MLP to $R^2=0.20$ (Table A.19, Appendix). With $\sim 250\text{K}$ -parameter models and ~ 62 samples, simpler architectures dominate—a finding consistent across the Exp-B iteration.

Table 10: Real-data regression (v3) and synthetic-to-real R^2 drop (v2→v3).

(a) Real-data metrics (97 days, 377 detections)					(b) R^2 drop (percentage points)			
Model	RMSE	R^2	F1	ms	Model	Syn.	Real	Δ
MLP	0.138	0.799	0.941	0.0	MLP	0.968	0.799	−17
LSTM	0.163	0.718	0.943	0.4	LSTM	0.951	0.718	−25
XGBoost	0.149	0.766	0.931	5.6	XGBoost	0.948	0.766	−19
Koopman-Det	0.157	0.739	0.924	0.1	Koopman-Det	0.968	0.739	−24
NeKo-PIGNN v2	0.152	0.756	0.926	0.3	NeKo-PIGNN v2	0.972	0.756	−22

6.9. Exp-B Rare-Event Detection and Three-Class Alerts (v5–v9)

v5–v7 reformulated the task as rare-event detection (positive rate 7.5%), then added persistence priors (semi-arid fires persist 2–7 days). NeKo-PIGNN achieved the highest Average Precision (0.353) in v5; ensemble $\times$ persistence reached 47.5% precision with 21 FP in v7 (Appendix A). **v8/v9** operationalize three response levels (NO/UNCERTAIN/YES), resolving the precision–recall trade-off documented in docs/EXPLICACAO_SISTEMA_DETECCAO.md.

6.10. Three-Level Operational Alert System (v8–v9)

Table 11 maps model classes to civil-defense actions and reports YES-class precision at operational probability thresholds. Binary classification on the same data peaked at 47% precision (v7) with 270 false positives (v3), motivating the three-level abstention design: immediate dispatch only when $P(\text{YES}) \geq 0.30$, monitoring for UNCERTAIN, and satellite-only coverage otherwise.

Table 11: Three-class alert mapping and YES-class precision at operational thresholds (20-day holdout).

(a) Alert levels				(b) YES thresholds				
Level	Class	Prec.	Rec.	Model	Prec.	Rec.	TP	FP
Alert	YES	82.1%	42%	XGB, $P \geq 0.30$	82.1%	41.8%	23	5
Watch	UNCERTAIN	—	+46%	XGB, $P \geq 0.50$	79.2%	34.6%	19	5
Safe	NO	—	12%	NeKo (YES)	91.7%	12.7%	11	1
Combined		—	88%	Ensemble, $P \geq 0.60$	88.9%	—	8	1
FP (alerts only)		5						

Table 12 summarizes how each Exp-B iteration reduced false positives while preserving coverage. Three-class abstention (v8/v9) adds +35 pp precision over binary MLP (v3); the persistence prior (v7) and weighted rare-event loss (v5–v6) provide intermediate gains documented in Appendix A. Figure 9 places all published operating points in precision–recall space: binary detection (v3–v7) is confined below the $F1=0.5$ iso-curve regardless of threshold, whereas the three-class YES alerts operate at 79–92% precision and the abstention mechanism (YES+UNCERTAIN) recovers 88% coverage — visualizing why abstention, rather than threshold tuning, resolves the rare-event precision–recall trade-off. The full per-class sklearn report and NeKo-PIGNN ablation are in the appendix (Tables A.23, A.24).




Figure 9: Published Exp-B operating points in precision–recall space (dotted lines: iso-F1 contours). Gray arrows trace the binary-detection iteration (v3→v7, Table 12); red markers show three-class YES alerts on the 20-day holdout (Table 11); stars show combined YES+UNCERTAIN coverage (Table 11) and the dry-season 2025 operational point (184 days, Table 14).

Table 12: Exp-B iteration history (v3–v9): precision–recall trade-off resolution.

Ver.	Approach	Prec.	Rec.	FP
v3	Binary MLP → alert	21%	96%	270
v5	NeKo + weighted loss	34%	27%	39
v6	XGBoost 300 trees	34%	59%	28
v7	+ Persistence prior	47%	21%	21
v8/v9	+ Three-class abstention	82%	88% [†]	5

Note: [†]Combined YES + UNCERTAIN coverage. NeKo YES-only: 91.7% prec., 12.7% rec., 1 FP.

6.11. Statistical Robustness (*Exp-ROBUST-001*)

Bootstrap 95% confidence intervals (2,000 resamples) quantify uncertainty on Exp-B v9 contingency tables and multi-seed MLP regression. XGBoost YES alert precision (82.1%) has CI [67.9%, 96.4%] (23 TP, 5 FP); combined YES+UNCERTAIN coverage 88% CI [84.6%, 91.2%]. Reproduced MLP regression yields RMSE 0.137 ± 0.003 and R^2 0.800 ± 0.010 [0.783, 0.811], consistent with Table 10. The alert-precision CIs resample the published v9 contingency tables directly (binomial bootstrap); the reproduction script (backend/experiments/statistical_robustness.py) additionally retrains a simplified in-script classifier used only for the paired tests below, whose more

conservative operating point is documented in the script and result files and should not be compared with the v9 pipeline.

Paired significance tests complement the bootstrap analysis (Table 13b). Wilcoxon signed-rank on per-seed MLP vs. XGBoost RMSE yields $p = 0.0625$; McNemar tests for three-class and YES-alert vs. persistence are not significant at $\alpha = 0.05$ on the 20-day holdout. We report these for transparency rather than as superiority claims.

Table 13: Bootstrap 95% CI and paired significance tests (Exp-ROBUST-001).

(a) Bootstrap CI				(b) Paired tests			
Metric	Point	Low	High	Test	p	Sig.	Notes
XGB YES prec.	82.1%	67.9%	96.4%	Wilcoxon	0.0625	No	$\Delta=-0.0099$
NeKo YES prec.	91.7%	75.0%	100.0%	RMSE			
Coverage	88.0%	84.6%	91.2%	McNemar	3- 0.1418	No	acc 73.2%
MLP RMSE	0.1372	0.1334	0.1429	class			
MLP R^2	0.800	0.783	0.811	McNemar YES	0.8714	No	prec 100%

6.12. Operational and Temporal Validation (Exp-ROBUST-003/006)

Table 14 positions the LangGraph pipeline against NASA FIRMS and INPE BDQueimadas, then tests alert generalization across temporal windows. Cold-start latency (52.4 s) is $\sim 309\times$ faster than the 3–6 h FIRMS publication window. Dry-season three-class alerts reach 84.2% precision and 91.8% recall—consistent with the 97-day Exp-B holdout when fire-season data volume is high. Mixed 12-month windows yield 31.6% precision, indicating season-specific calibration. Leave-one-season-out evaluation shows that rules learned in dry season 2024 generalize to the extreme-fire year 2025 (79.0% precision, 85.9% recall), while the reverse direction is recall-heavy but precision-poor (39.2%), supporting season-aware deployment.

Table 14: Operational baselines and extended temporal validation (XGBoost 3-class, YES at $P \geq 0.30$).

(a) System comparison				(b) Temporal windows				
System	Lat.	Prec.	Rec.	Window	Days	Prec.	Rec.	FP
NASA FIRMS	3–6 h	80.0%	15.7%	Short sub-set	117	34.4%	38.9%	40
INPE BD-Queimadas	3–6 h	—	—	Dry 2024	141	44.9%	66.4%	114
LangGraph (ours)	52 s	84.2%	91.8%	Dry 2025	184	84.2%	91.8%	75
<i>Note:</i> FIRMS→INPE match 25.8% (1 km, 8-day sample). Cost: \$20/mo.				Full mo	12	31.6%	74.1%	130
				LOO 24→25	500	79.0%	85.9%	379
				LOO 25→24	500	39.2%	98.2%	595

6.13. External-State Generalization (Exp-ROBUST-005)

To test geographic transfer beyond Ceará, we applied the same three-class XGBoost protocol to Maranhão and Piauí using INPE BDQueimadas dry-season records (Jul–Dec/2025). Ceará retains the best precision–recall balance (84.2% / 91.8%); neighboring states show higher fire volume and recall (>96%) but lower precision (~58–60%), consistent with class imbalance and semi-arid persistence patterns (Table 15).

Table 15: External-state validation: XGBoost 3-class YES alert at $P \geq 0.30$ (dry season Jul–Dec/2025).

State	Days	Hotspots	Precision	Recall
Ceará	184	24,573	84.2%	91.8%
Maranhão	184	233,854	57.5%	99.5%
Piauí	183	145,220	60.1%	96.4%

Figure 10 visualizes the precision–recall trade-off across states. These results support a Northeast pilot deployment model: calibrate thresholds per state rather than transferring Ceará parameters directly to higher-volume neighbors.




Figure 10: Dry-season 2025 alert precision and recall: Ceará vs. Maranhão and Piauí (INPE BDQueimadas).

6.14. Comparison with Published Systems

Table 16 situates the platform relative to operational services (INPE, FIRMS), research Digital Twins (IVSR, FIRE-VLM), physics simulators (FlamMap/FARSITE), and published ML baselines. The distinguishing combination is sub-minute LangGraph orchestration, three-class alert precision (82–92%), and open-source deployment on a single EC2 instance—at the cost of state-specific calibration and limited GOES-16 production integration.

Table 16: Operational platforms and published wildfire detection baselines.

(a) Operational systems						(b) ML / index baselines			
System	Lat.	Prec.	Ag.	Scope		System	Method	Prec.	Metric
INPE	BD-	3–6 h	ref.	No	Brazil	Canadian FWI	Physical	—	AUC 0.69
Queimadas						Regional	Tailored	—	AUC 0.79
NASA		3–6 h	~80%	No	Global	FWI+GA			
FIRMS						Decision Tree	ML/country	—	AUC 0.86
FIRE-VLM	sim.	—	RL	Research		FC	Autoen-	—	F1 0.74
IVSR	reduced	—	Yes	Multi-sensor		coder	Anomaly		
This platform	52 s	84%	Yes	CE+NE		This work	GNN + Koopman	82–92%	F1 0.54

Note: Macro F1 from Table A.23.

6.15. Predictive Model Validation: Summary

Sections 6.7–6.10 report the complete Exp-B iteration (v1–v9): synthetic baselines where NeKo-PIGNN reaches $R^2=0.972$; real-data regression where MLP leads RMSE with 67 training days; rare-event binary detection (v5–v7) where NeKo-PIGNN achieves the highest Average Precision (0.353); and the final three-class alert system (v8/v9) with 82–92% YES precision and 88% combined coverage. Binary regression F1 (~ 0.94 , Table 10) and operational YES precision (82–92%) measure *different tasks* and must not be conflated.

6.16. Operational Deployment Metrics

Table 17 summarizes 24 months of continuous deployment: >99% backend availability, \$20 USD/month infrastructure cost, and 100% free/open data sources.

Table 17: Operational deployment metrics (Jul/2024–Jun/2026).

Metric	Value
Operation time	24 months
FIRMS collections	540+
FIRMS hotspots detected	38,400+
INPE hotspots (reference period)	111,639
Backend availability	>99 %
Monthly infrastructure cost	\$20 USD
Paid data sources	0 %

7. Discussion

The results show that open satellite and climate streams can support timely ecological fire monitoring in the Caatinga without institutional databases or paid APIs. Relative to raw FIRMS (3–6 h latency; 15.7% INPE spatial recall on the bundled sample), the LangGraph stack trades modest cost (\$20/month) for sub-minute cold starts and 91.8% dry-season alert recall (Table 14). For ecosystem managers, three-class alerts (YES/UNCERTAIN/NO) reduce false alarms while

preserving coverage of fire-prone municipalities — a practical bridge between remote sensing and sustainable fire management under climate-driven drought.

7.1. Ecological Implications for the Semi-Arid Northeast

Ceará’s fire regime is tightly coupled to rainfall deficit, fuel curing, and land-use pressure in a biodiversity hotspot with limited fire-suppression infrastructure. High dry-season precision/recall (84.2%/91.8%) indicates that open VIIRS/MODIS products, when fused with Open-Meteo drought indicators, can prioritize municipal attention during the peak ecological risk window. Lower precision in wetter or mixed seasons (Table 14) and in high-volume neighboring states (Section 6.13) underscores that ecological transferability requires per-biome threshold calibration rather than a single national cutoff.

7.2. NeKo-PIGNN and Operational Scope

NeKo-PIGNN remains an *offline* research module: under 67-day real windows, MLP/XGBoost achieve lower RMSE, while NeKo contributes synthetic validation ($R^2 = 0.972$), conservative YES precision (91.7%, 1 FP), and physics-regularized spatial propagation for future ecological forecasting. Production alerts rely on the three-class LangGraph stack, not NeKo RMSE.

7.3. Auditability, Coverage, and Limitations

ReAct traces make each alert auditable for environmental agencies. Multisensor VIIRS/MODIS coverage (6–8 daily passes) reduces single-sensor blind spots, but wet-season cloud cover and optional GOES-16 ($F1 \leq 0.05$ vs. INPE at grid scale) remain constraints. Additional limitations: discontinuous INPE API access during parts of the study; NeKo real-data scarcity (67 days; YES recall 12.7%); RAG recall@5 = 98% bilingual vs. 74% English-only; no formal SUS usability trial yet (informal stakeholder feedback only). Risk-index calibration against an INPE snapshot reached $F1=0.936$ (Section 6.6). No institutional or academic system surveyed here simultaneously offers sub-60 s open-source response, ReAct audit trails, RAG, and Caatinga dry-season validation.

8. Conclusion and Future Work

We presented an open-source ecological monitoring platform for wildfires in Ceará’s Caatinga, integrating open satellite and climate data with agentic AI and a physics-informed research module for risk-aware decision support.

The main contributions are: (1) a three-class alert framework with **84.2% precision and 91.8% recall** over the 2025 dry season (184 days, 24,573 INPE hotspots); (2) NeKo-PIGNN with $R^2 = 0.972$ on synthetic dynamics and transparent real-data limits; (3) bootstrap confidence intervals and multi-seed checks; and (4) an MIT-licensed pipeline on a low-cost EC2 instance.

Three-class abstention identifies 88% of fire events at 82% YES precision on short windows. Future work includes production GOES-16 fusion, MapBiomas cross-validation, Northeast expansion with per-state calibration, sequence predictors, messaging for municipal agencies, and a formal SUS usability study.

Acknowledgments

The authors thank the University of Fortaleza (UNIFOR) for institutional support, NASA FIRMS and NOAA for public access to satellite data, and Open-Meteo for the free weather API.

Software and data availability

Name of software: Ceará Wildfire Monitoring Platform

Developers: Nauber Gois et al. (Universidade de Fortaleza — UNIFOR)

Contact: naubergois@gmail.com

Year first available: 2024

Program language: Python 3.12 (backend, models), TypeScript/React 19 (frontend)

Hardware requirements: Single cloud instance (2 vCPUs, 4 GB RAM; tested on AWS EC2 t2.medium)

Software requirements: Linux, Python 3.12, Node.js 20, nginx; optional Docker Compose (PostgreSQL/PostGIS, Redis)

Cost: Free and open source (MIT license)

Availability: <https://github.com/naubergois/ceara-queimadas>

All experiment scripts, logs, and result JSON files are included in the repository. Key reproducibility paths:

- Exp-ROBUST-001: `backend/experiments/statistical_robustness.py`
- Exp-ROBUST-002: `src/pixel_inpe_eval.py`
- Exp-ROBUST-003: `backend/experiments/extended_temporal_benchmark.py`
- Exp-ROBUST-004: `backend/experiments/rag_benchmark.py`
- Exp-ROBUST-005: `backend/experiments/external_state_benchmark.py`
- Exp-ROBUST-006: `backend/experiments/operational_baseline_benchmark.py`
- PR operating-point map (Fig. 9): `backend/experiments/pr_curve_figure.py`
- INPE CSVs: `scripts/reproduce_inpe_data.sh` → `data/inpe_focos_ce/`
- Supplementary tables: `submission/supplementary.tex`
- Reproducibility archive: GitHub release v1.0.0 - ems - submission at <https://github.com/naubergois/ceara-queimadas/releases> (Zenodo DOI optional — see `submission/ZENODO.md`)

CRediT authorship contribution statement

Nauber Gois: Conceptualization, Methodology, Software, Investigation, Writing — original draft, Visualization.

Alexandre Lobo: Conceptualization, Methodology, Supervision, Validation, Data Curation, Writing — review and editing.

Vladia Celia Monteiro Pinheiro: Conceptualization, Methodology, Supervision, Writing — review and editing.

Daniel Valente de Macedo: Conceptualization, Methodology, Writing — review and editing.

Adriano Bessa Albuquerque: Conceptualization, Methodology, Writing — review and editing.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Declaration of generative AI and AI-assisted technologies in the writing process

During the preparation of this work the authors used AI-assisted writing tools for language refinement, literature search, and L^AT_EX formatting. After using these tools, the authors reviewed and edited the content as needed and take full responsibility for the content of the published article.

Appendix A. Supplementary Exp-B Tables

Detailed iteration tables for Exp-B v1, v4–v7, the full three-class sklearn report, and the municipal hotspot bar chart referenced in Section 6.2.

Figure A.11 shows municipal hotspot counts for the Exp-B real-data subset. Tables A.18–A.22 document intermediate Exp-B iterations referenced in Section 6.7; Table A.23 gives the full sklearn report; Table A.24 reports NeKo-PIGNN component ablation on synthetic Rothermel propagation data.




Figure A.11: Hotspots per monitored municipality (Exp-B real-data subset, Mar–Jun/2026).

Table A.18 records the initial synthetic benchmark where KL-regularized Koopman-VAE underperformed shallow baselines. Table A.19 shows that feature expansion degraded MLP performance under data scarcity (62–67 training samples).

Tables A.20–A.22 trace the rare-event detection iteration: v5 introduces Average Precision optimization; v6 adds focal loss and conservative XGBoost; v7 incorporates the persistence prior that raised precision to 47.5% with only 21 false positives.

Table A.23 reports per-class precision/recall/F1 from argmax classification; operational YES alerts use $P(\text{YES}) \geq 0.30$ instead (Section 6.10). Table A.24 confirms synergistic Koopman-PI-GNN coupling (full F1 = 0.914 vs. Koopman-only 0.820).

Table A.18: Exp-B v1: synthetic regression (200 steps, 20 nodes).

Model	RMSE↓	R ² ↑	F1	Recall
MLP	0.083	0.932	0.957	0.941
XGBoost	0.085	0.929	0.958	0.945
Koopman-VAE (ours)	0.260	0.343	0.828	0.989
NeKo-PIGNN (ours)	0.272	0.276	0.830	0.989

Table A.19: Exp-B v4: real-data regression with feature-engineering variants.

Model	RMSE↓	R ² ↑	F1	Precision
MLP plain	0.143	0.782	0.936	0.911
XGBoost + FeatEng	0.149	0.767	0.935	0.902
Ensemble (NeKo+XGB)	0.149	0.766	0.930	0.896
NeKo-PIGNN v3	0.173	0.686	0.906	0.907
MLP + FeatEng + LB7	0.276	0.197	0.830	0.936

Table A.20: Exp-B v5: rare-event detection — Average Precision and threshold metrics.

Model	AP↑	F1@best	Prec.	Recall	FP@0.3
MLP	0.291	0.348	0.212	0.978	270
NeKo-PIGNN v5	0.353	0.323	0.385	0.278	39
XGBoost	0.304	0.295	0.349	0.256	43
Ensemble (NeKo+MLP)	0.325	0.345	0.455	0.278	—

Table A.21: Exp-B v6: precision–recall operating modes.

Model	AP	F1@best	Prec.	Recall	FP@0.3
XGBoost (300 trees)	0.322	0.427	0.335	0.589	28
Ensemble (NeKo+XGB+MLP)	0.315	0.353	0.217	0.956	67
NeKo-PIGNN (5-seed avg.)	0.280	0.324	0.233	0.533	81
MLP (5-seed)	0.276	0.350	0.213	0.989	282

Table A.22: Exp-B v7: persistence prior combined with ML ensemble.

Model	Precision	Recall	F1	FP
Persistence only	45.1%	48.7%	0.469	45
Ensemble × Persistence	47.5%	21.1%	0.292	21
NeKo × Persistence	40.7%	24.4%	0.306	32
XGBoost raw	33.5%	58.9%	0.427	105

Table A.23: Full three-class classification report (Exp-B v9, XGBoost, 420 samples).

Class	Precision	Recall	F1	Support
NO	0.693	0.973	0.810	186
UNCERTAIN	0.669	0.674	0.671	144
YES	0.500	0.078	0.135	90
Macro avg	0.621	0.575	0.539	420
Weighted avg	0.644	0.679	0.618	420

Note: Operational YES alerts use $P(\text{YES}) \geq 0.3$ (82.1% precision), not argmax class.

Table A.24: NeKo-PIGNN component ablation on synthetic Rothermel propagation data.

Model	F1	IoU	RMSE	R ²
NeKo-PIGNN (full)	0.9140	0.8418	0.1083	0.8102
Koopman-only	0.8200	0.6949	0.1531	0.6215
PI-GNN-only	0.7900	0.6529	0.1682	0.5394
Pure GNN	0.3678	0.2255	0.2851	−0.32
CNN U-Net	0.4599	0.2990	0.2514	−0.03

References

- [1] V. R. Pivello, I. Vieira, A. V. Christianini, D. B. Ribeiro, L. da Silva Menezes, C. N. Berlinck, F. P. L. Melo, J. A. Marengo, C. G. Tornquist, W. M. Tomas, G. E. Overbeck, Understanding Brazil’s catastrophic fires: Causes, consequences and policy needed to prevent future tragedies, *Perspectives in Ecology and Conservation* 19 (3) (2021) 233–255. doi:10.1016/j.pecon.2021.06.005.
- [2] J. Silva, Impacts of wildfires on public health in the Brazilian semi-arid region, Master’s thesis, Universidade Federal do Ceará (2021).
- [3] FUNCEME, Drought monitoring bulletin — Ceará 2024, Tech. rep., FUNCEME (2024).
- [4] A. A. C. Alencar, V. L. S. Arruda, W. V. d. Silva, D. E. Conciani, D. P. Costa, N. Crusco, S. G. Duverger, N. C. Ferreira, W. Franca-Rocha, H. Hasenack, et al., Long-Term Landsat-Based Monthly Burned Area Dataset for the Brazilian Biomes Using Deep Learning, *Remote Sensing* 14 (11) (2022) 2510. doi:10.3390/rs14112510.
- [5] NASA FIRMS, Fire Information for Resource Management System — User Guide (2024). URL <https://firms.modaps.eosdis.nasa.gov/>
- [6] S. Yao, J. Zhao, D. Yu, et al., ReAct: Synergizing Reasoning and Acting in Language Models, in: *International Conference on Learning Representations (ICLR)*, 2023.
- [7] LangChain, LangGraph Documentation (2024). URL <https://langchain-ai.github.io/langgraph/>

- [8] P. Bauer, B. Stevens, W. Hazeleger, A digital twin of Earth for the green transition, *Nature Climate Change* 11 (2) (2021) 80–83. doi:10.1038/s41558-021-00986-y.
- [9] M. Grieves, *Digital Twin: Manufacturing Excellence through Virtual Factory Replication*, White paper, Florida Institute of Technology (2014).
- [10] F. Tao, H. Zhang, A. Liu, A. Y. C. Nee, Digital twin in industry: State-of-the-art, *IEEE Transactions on Industrial Informatics* 15 (4) (2019) 2405–2415. doi:10.1109/TII.2018.2873186.
- [11] B. R. Barricelli, E. Casiraghi, D. Fogli, A survey on digital twin: Definitions, characteristics, applications, and design implications, *IEEE Access* 7 (2019) 167653–167671. doi:10.1109/ACCESS.2019.2953499.
- [12] A. Fuller, Z. Fan, C. Day, C. Barlow, Digital twin: Enabling technologies, challenges and open research, *IEEE Access* 8 (2020) 108952–108971. doi:10.1109/ACCESS.2020.2998358.
- [13] S. Nativi, P. Mazzetti, M. Craglia, Digital Ecosystems for Developing Digital Twins of the Earth: The Destination Earth Case, *Remote Sensing* 13 (11) (2021) 2119. doi:10.3390/rs13112119.
- [14] P. Voosen, Europe builds 'digital twin' of earth to hone climate forecasts, *Science* 370 (6512) (2020) 16–17. doi:10.1126/science.370.6512.16.
- [15] M. Morsali, S. H. Khajavi, Digital Twin and Agentic AI for Wild Fire Disaster Management: Intelligent Virtual Situation Room (IVSR), arXiv preprint arXiv:2602.08949 (2026). arXiv: 2602.08949.
URL <https://arxiv.org/abs/2602.08949>
- [16] C. Webb, M. Habibpour, M. H. Raha, A. R. Tavakkoli, J. Coen, F. Afghah, FIRE-VLM: A Vision-Language-Driven Reinforcement Learning Framework for UAV Wildfire Tracking in a Physics-Grounded Fire Digital Twin, arXiv preprint arXiv:2601.03449 (2026). arXiv: 2601.03449.
URL <https://arxiv.org/abs/2601.03449>
- [17] M. H. Raha, A. R. Tavakkoli, C. Webb, M. Habibpour, J. Coen, E. Rowell, F. Afghah, FIRETWIN: Digital Twin Advancing Multi-Modal Sensing, Interactive Analytics for Tactical Wildfire Response, in: *2025 IEEE 30th International Workshop on Computer Aided Modeling and Design of Communication Links and Networks (CAMAD)*, 2025, pp. 1–6. arXiv: 2510.18879, doi:10.1109/camad67323.2025.11229918.
URL <https://arxiv.org/abs/2510.18879>
- [18] Z. Zhou, X. Wang, Y. Yan, L. Mijiddorj, Y. Ding, T. Beringer, P. M. Khiabani, W. G. Jentner, X.-M. Hu, C. Wang, B. M. Carroll, M. Xue, D. Ebert, J. Li, AIMNET: An IoT-Empowered Digital Twin for Continuous Gas Emission Monitoring and Early Hazard Detection, arXiv preprint arXiv:2512.06148 (2025). arXiv: 2512.06148.
URL <https://arxiv.org/abs/2512.06148>

- [19] I. Esparza, A. Battal, A. Mostafavi, GraphFire-X: Physics-Informed Graph Attention Networks and Structural Gradient Boosting for Building-Scale Wildfire Preparedness at the Wildland-Urban Interface, arXiv preprint arXiv:2512.20813 (2025). arXiv: 2512. 20813.
URL <https://arxiv.org/abs/2512.20813>
- [20] D. Michail, I. Davalas, K. Chafis, L. I. Panagiotou, I. Prapas, N. I. Bountos, S. Kondylatos, I. Papoutsis, FireCastNet: Earth-as-a-Graph for Seasonal Fire Prediction, Scientific Reports (2025). arXiv:2502.01550, doi:10.1038/s41598-025-30645-7.
URL <https://arxiv.org/abs/2502.01550>
- [21] L. Giglio, W. Schroeder, C. O. Justice, The collection 6 modis active fire detection algorithm and fire products, Remote Sensing of Environment 178 (2016) 31–41. doi:10.1016/j.rse.2016.02.054.
- [22] W. Schroeder, P. Oliva, L. Giglio, I. A. Csizsar, The new viirs 375m active fire detection data product: Algorithm description and initial assessment, Remote Sensing of Environment 143 (2014) 85–96. doi:10.1016/j.rse.2013.12.008.
- [23] INPE, Wildfire Database — Methodology and Data (2023).
URL <https://queimadas.dgi.inpe.br/>
- [24] M. J. Wooster, G. J. Roberts, L. Giglio, et al., Satellite remote sensing of active fires: History and current status, applications and future requirements, Remote Sensing of Environment 267 (2021) 112694. doi:10.1016/j.rse.2021.112694.
- [25] NOAA, GOES-16 ABI L2+ Fire Detection and Characterization, Tech. rep., NOAA (2022).
- [26] MapBiomas, MapBiomas Fogo, accessed 2025 (2025).
URL <https://mapbiomas.org/>
- [27] J. Barlow, E. Berenguer, R. Carmenta, F. França, Clarifying Amazonia’s burning crisis, Global Change Biology 26 (2) (2020) 319–321. doi:10.1111/gcb.14872.
- [28] P. Jain, S. C. P. Coogan, S. G. Subramanian, M. Crowley, S. Taylor, M. D. Flannigan, A review of machine learning applications in wildfire science and management, Environmental Reviews 28 (4) (2020) 478–505. doi:10.1139/er-2020-0019.
- [29] B. O. Koopman, Hamiltonian Systems and Transformation in Hilbert Space, Proceedings of the National Academy of Sciences 17 (5) (1931) 315–318. doi:10.1073/pnas.17.5.315.
- [30] I. Mezić, Analysis of fluid flows via spectral properties of the koopman operator, Annual Review of Fluid Mechanics 45 (2013) 357–378. doi:10.1146/annurev-fluid-011212-140652.
- [31] P. J. Schmid, Dynamic mode decomposition of numerical and experimental data, Journal of Fluid Mechanics 656 (2010) 5–28. doi:10.1017/S0022112010001217.
- [32] S. L. Brunton, M. Budišić, E. Kaiser, J. N. Kutz, Modern Koopman Theory for Dynamical Systems, SIAM Review 64 (1) (2022) 229–340. doi:10.1137/21M1401243.

- [33] B. Lusch, J. N. Kutz, S. L. Brunton, Deep learning for universal linear embeddings of nonlinear dynamics, *Nature Communications* 9 (2018) 4950. doi:10.1038/s41467-018-07210-0.
- [34] M. Raissi, P. Perdikaris, G. E. Karniadakis, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, *Journal of Computational Physics* 378 (2019) 686–707. doi:10.1016/j.jcp.2018.10.045.
- [35] G. E. Karniadakis, I. G. Kevrekidis, L. Lu, P. Perdikaris, S. Wang, L. Yang, Physics-informed machine learning, *Nature Reviews Physics* 3 (6) (2021) 422–440. doi:10.1038/s42254-021-00314-5.
- [36] R. C. Rothermel, A mathematical model for predicting fire spread in wildland fuels, Tech. Rep. INT-115, USDA Forest Service (1972).
- [37] J. H. Scott, R. E. Burgan, Standard fire behavior fuel models: a comprehensive set for use with Rothermel’s surface fire spread model, Tech. Rep. RMRS-GTR-153, USDA Forest Service (2005).
- [38] P. Lewis, E. Perez, A. Piktus, et al., Retrieval-augmented generation for knowledge-intensive NLP tasks, in: *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [39] J. Johnson, M. Douze, H. Jégou, Billion-scale similarity search with GPUs, *IEEE Transactions on Big Data* 7 (3) (2021) 535–547. doi:10.1109/TBDATA.2019.2921572.
- [40] Y. Gao, Y. Xiong, X. Gao, K. Jia, J. Pan, Y. Bi, Y. Dai, J. Sun, Q. Guo, M. Wang, H. Wang, Retrieval-augmented generation for large language models: A survey, *arXiv preprint arXiv:2312.10997* (2024). arXiv:2312.10997.
URL <https://arxiv.org/abs/2312.10997>
- [41] BAAI, BGE Small Embedding Model v1.5 (2024).
URL <https://huggingface.co/BAAI/bge-small-en-v1.5>
- [42] C. W. Rowley, I. Mezić, S. Bagheri, P. Schlatter, D. S. Henningson, Spectral analysis of nonlinear flows, *Journal of Fluid Mechanics* 641 (2009) 115–127. doi:10.1017/S0022112009992059.
- [43] N. Takeishi, Y. Kawahara, T. Yairi, Learning koopman invariant subspaces for dynamic mode decomposition, in: *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 30, 2017.
- [44] M. O. Williams, I. G. Kevrekidis, C. W. Rowley, A Data-Driven Approximation of the Koopman Operator: Extending Dynamic Mode Decomposition, *Journal of Nonlinear Science* 25 (6) (2015) 1307–1346. doi:10.1007/s00332-015-9258-5.
- [45] L. Lu, X. Meng, Z. Mao, G. E. Karniadakis, Deepxde: A deep learning library for solving differential equations, *SIAM Review* 63 (1) (2021) 208–228. doi:10.1137/19M1274067.

- [46] S. Cuomo, V. S. D. Cola, F. Giampaolo, G. Rozza, M. Raissi, F. Piccialli, Scientific machine learning through physics-informed neural networks: Where we are and what's next, *Journal of Scientific Computing* 92 (3) (2022) 88. doi:10.1007/s10915-022-01939-z.
- [47] M. Fey, J. E. Lenssen, Fast graph representation learning with PyTorch Geometric, in: *ICLR Workshop on Representation Learning on Graphs and Manifolds*, 2019.
- [48] Y. Bengio, J. Louradour, R. Collobert, J. Weston, Curriculum learning, in: *Proceedings of the 26th International Conference on Machine Learning (ICML)*, 2009, pp. 41–48. doi:10.1145/1553374.1553380.